

```
In [1]: import os
import numpy as np
import torch as T
import torch.nn as nn
import torch.optim as optim
from torch.distributions.normal import Normal
from torch.nn import functional as F
import matplotlib.pyplot as plt
import gym
```

```
D:\miniconda3\envs\cmpe260\lib\site-packages\tqdm\auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
```

```
In [2]: import a3_gym_env
gym.make('Pendulum-v260')
```

```
Out[2]: <OrderEnforcing<PassiveEnvChecker<CustomPendulumEnv<Pendulum-v260>>>>
```

```
In [3]: class ReplayBuffer:
    def __init__(self):
        self.states = []
        self.probs = []
        self.vals = []
        self.actions = []
        self.rewards = []
        self.dones = []
        self.h_prevs = []

    def store_memory(self, state, action, probs, vals, reward, done, h_prevs):
        self.states.append(state)
        self.actions.append(action)
        self.probs.append(probs)
        self.vals.append(vals)
        self.rewards.append(reward)
        self.dones.append(done)
        self.h_prevs.append(h_prevs)
```

```
In [4]: class ActorNetwork(nn.Module):
    def __init__(self, input_dims, output_dims, alpha,
                  hidden_layers = 2, hidden_dims=8, N=1):
        super().__init__()
        # actor network to estimate mean of the gaussian distribution
        self.hidden_dims = hidden_dims
        self.lin1 = nn.Linear(input_dims, hidden_dims)
        self.lstm = nn.LSTM(hidden_dims, hidden_dims)
        self.lin2 = nn.Linear(hidden_dims, output_dims)

        # use constant standard deviation
        log_std = -0.9 * np.ones(output_dims, dtype=np.float32)
        self.log_std = T.nn.Parameter(T.as_tensor(log_std))

        self.optimizer = optim.Adam(self.parameters(), lr=alpha/N, eps=1e-05)
        self.device = T.device('cuda:0' if T.cuda.is_available() else 'cpu')
        self.to(self.device)

    def forward(self, state):
        obs, h_prev = state
        m = F.relu(self.lin1(obs))
        m = m.view(-1, 1, self.hidden_dims)
        m, h_next = self.lstm(m, h_prev)
        m = self.lin2(m)
#         m = outputs[-1]
        # return normal distribution
        dist = Normal(m, T.exp(self.log_std))
        return dist, h_next
```

```
In [5]: class CriticNetwork(nn.Module):
    def __init__(self, input_dims, alpha, hidden_layers = 2, hidden_dims=8, N=1):
        super().__init__()
        # critic network for estimating value
        self.hidden_dims = hidden_dims
        self.lin1 = nn.Linear(input_dims, hidden_dims)
        self.lstm = nn.LSTM(hidden_dims, hidden_dims)
        self.lin2 = nn.Linear(hidden_dims, 1)

        self.optimizer = optim.Adam(self.parameters(), lr=alpha/N, eps=1e-05)
        self.device = T.device('cuda:0' if T.cuda.is_available() else 'cpu')
        self.to(self.device)

    def forward(self, state):
        # return value
        obs, h_prev = state
#         import pdb; pdb.set_trace()
        value = F.relu(self.lin1(obs))
        value = value.view(-1, 1, self.hidden_dims)
        value, _ = self.lstm(value, h_prev)
        value = self.lin2(value)
        return value
```

```
In [6]: def choose_action_value(observation, h_prev, actor, critic):
    # convert state to tensor
    state = T.tensor([observation], dtype=T.float).to(actor.device)
    # get gaussian distribution for state
    dist, h_next = actor((state, h_prev))
    # get value for current state
    value = critic((state, h_prev))
    # sample action
    action = dist.sample()

    # get log prob for sampled action (to be used for ratio)
    probs = T.squeeze(dist.log_prob(action)).item()
    action = T.squeeze(action).item()
    value = T.squeeze(value).item()
    # return action, log prob for sampled action and value for current state
    return action, probs, value, h_next
```

```
In [7]: def get_advantages(gamma, lambda, values, done, rewards,
    next_value, next_done, use_gae):
    num_steps = len(rewards)
    if use_gae:
        advantages = np.zeros_like(rewards)
        last_gae_lam = 0
        for t in reversed(range(num_steps)):
            if t == num_steps - 1:
                next_non_terminal = 1.0 - next_done
                next_values = next_value
            else:
                next_non_terminal = 1.0 - done[t + 1]
                next_values = values[t + 1]
            delta = rewards[t] + gamma * next_values * next_non_terminal - values[t]
            advantages[t] = last_gae_lam = delta + gamma * lambda * next_non_terminal * last_gae_lam

        returns = advantages + values
    else:
        returns = np.zeros_like(rewards)
        for t in reversed(range(num_steps)):
            if t == num_steps - 1:
                next_non_terminal = 1.0 - next_done
                next_return = next_value
            else:
                next_non_terminal = 1.0 - done[t + 1]
                next_return = returns[t + 1]
            returns[t] = rewards[t] + gamma * next_non_terminal * next_return
        advantages = returns - values
    return returns, advantages
```

```
In [8]: def combine_trajectories(trajectories):
c_states = []
c_actions = []
c_rewards = []
c_logprobs = []
c_values = []
c_dones = []
c_returns = []
c_advantages = []
c_h_prevs = []
for k, v in trajectories.items():
    buf, ret, adv = v
    c_states.extend(buf.states)
    c_actions.extend(buf.actions)
    c_rewards.extend(buf.rewards)
    c_logprobs.extend(buf.probs)
    c_values.extend(buf.vals)
    c_dones.extend(buf.dones)
    c_returns.extend(ret)
    c_advantages.extend(adv)
    c_h_prevs.extend(buf.h_prevs)
return {
    'states': c_states,
    'actions': c_actions,
    'rewards': c_rewards,
    'logprobs': c_logprobs,
    'values': c_values,
    'dones': c_dones,
    'returns': c_returns,
    'advantages': c_advantages,
    'h_prevs': c_h_prevs,
}
```

```

In [9]: def ppo(N = 1, M = 50, max_trajectory_len = 200, batch_size = 25,
           alpha = 0.0001, hidden_layers = 2, hidden_dims = 8, gamma = 0.99,
           lambda = 0.95, clip_value = 0.2, use_gae=True, use_clip=True):
    # N = 1 # number of times to collect new trajectories and update actor
    # M = 50 # num of trajectories
    # max_trajectory_len = 200 # trajectory length
    # batch_size = 25 # size for minibatch
    # n_epochs = 30 # number of epochs to optimize loss
    # alpha = 0.00005
    # hidden_layers = 3
    # hidden_dims = 10
    # gamma = 0.8
    # lambda = 0.95
    # clip_value = 0.2
    # use_gae = True
    # use_clip = True

    env = gym.make('Pendulum-v260')
    # input_dim = env.observation_space.shape
    input_dim = 1
    actor = ActorNetwork(input_dims=input_dim, output_dims=1, alpha=alpha,
                        hidden_layers=hidden_layers, hidden_dims=hidden_dims)
    critic = CriticNetwork(input_dims=input_dim, alpha=alpha,
                          hidden_layers=hidden_layers, hidden_dims=hidden_dims)

    trajectories = {}
    total_loss_list = []
    loss_list = []
    reward_list = []
    observation_list = []
    for n in range(N):
        # collect M trajectories, their returns, and activations
        for i in range(M):
            curr_trajectory = ReplayBuffer()
            # hprev and cell state
            h_prev = (T.zeros([1, 1, hidden_dims], dtype=T.float).to(actor.device),
                     T.zeros([1, 1, hidden_dims], dtype=T.float).to(actor.device))
            observation = env.reset()
            # only use angular velocity
            observation = observation[-1]
            done = False
            steps = 0
            while steps < max_trajectory_len:
                action, prob, val, h_next = choose_action_value(
                    observation, h_prev, actor, critic)
                next_observation, reward, done, info = env.step([action])
                steps += 1
                curr_trajectory.store_memory(
                    observation, action, prob, val, reward, done, h_prev)
                observation = next_observation[-1]
                h_prev = h_next
                if done:
                    break
            # get GAE advantage
            # get value of next_observation
            next_value = critic((T.tensor([observation]).to(actor.device), h_prev))
            # calculate return and advantage
            returns, advantages = get_advantages(
                gamma, lambda, curr_trajectory.vals, curr_trajectory.dones,
                curr_trajectory.rewards, next_value, done, use_gae)
            trajectories[i] = (curr_trajectory, returns, advantages)

    # combine_trajectories
    c_trajectories = combine_trajectories(trajectories)
    total_steps_taken = len(c_trajectories['states'])
    indices = np.arange(total_steps_taken)
    np.random.shuffle(indices)

```

```

np.random.shuffle(indicies)

# optimize loss
for epoch in range(n_epochs):
    # get batches
    for start in range(0, total_steps_taken, batch_size):
        end = start + batch_size
        # batch indicies
        b_indicies = indicies[start:end]
        new_value = []
        distr = []
        for index in b_indicies:
            state = T.tensor([c_trajectories['states'][index]]).to(actor.device)
            hprev = c_trajectories['h_prevs'][index]
            new_value.append(critic((state, hprev)))
            distr.append(actor((state, hprev)))

        new_value = T.tensor(new_value).to(actor.device)
        #
        # states = T.tensor(
        #     np.array(c_trajectories['states'])[b_indicies],
        #     dtype=T.float).to(actor.device)
        #
        # import pdb; pdb.set_trace()
        # h_prevs = c_trajectories['h_prevs'][b_indicies]
        # new_value = critic((states, h_prevs))
        # distr = actor((states, h_prevs))

        old_log_prob = T.tensor(
            np.array(c_trajectories['logprobs'])[b_indicies]).to(actor.device)
        actions = T.tensor(
            np.array(c_trajectories['actions'])[b_indicies],
            dtype=T.float).to(actor.device)
        new_log_prob = T.tensor([d[0].log_prob(actions[i]) for i, d in enumerate(distr)
            dtype=T.float]).to(actor.device)

        # Probability ratio
        log_ratio = new_log_prob - old_log_prob
        ratio = log_ratio.exp()

        b_returns = T.tensor(
            np.array(c_trajectories['returns'])[b_indicies]).to(actor.device)
        b_advantages = T.tensor(
            np.array(c_trajectories['advantages'])[b_indicies]).to(actor.device)
        #
        b_advantages = (b_advantages - b_advantages.mean()) / (b_advantages.std() + 1e-8)
        # Clipping
        weighted_probs = b_advantages * ratio
        if use_clip:
            weighted_clipped_probs = T.clamp(ratio, 1-clip_value,
                1+clip_value) * b_advantages
            actor_loss = -T.min(
                weighted_probs, weighted_clipped_probs).mean()
        else:
            weighted_probs = b_advantages * ratio
            actor_loss = -(weighted_probs).mean()
        critic_loss = (b_returns - new_value)**2
        critic_loss = critic_loss.mean()

        total_loss = actor_loss + 0.5 * critic_loss
        total_loss.requires_grad = True
        print("n: {}, epoch: {}, minibatch no.: {}, ".format(
            n, epoch, start // batch_size))
        print("total_loss: {}, actor_loss: {}, critic_loss: {}".format(
            total_loss, actor_loss, critic_loss))

        total_loss_list.append(total_loss)
        actor.optimizer.zero_grad()
        critic.optimizer.zero_grad()
        total_loss.backward()

```

```

        total_loss.backward()
        actor.optimizer.step()
        critic.optimizer.step()
#         actor.scheduler.step(actor_loss)
#         critic.scheduler.step(critic_loss)

# get a single trajectory from actor for plotting
loss_list.append(total_loss)
reward_list.append(b_returns.mean())

observation = env.reset()
h_prev = (T.zeros([1, 1, hidden_dims], dtype=T.float).to(actor.device),
          T.zeros([1, 1, hidden_dims], dtype=T.float).to(actor.device))
obs = observation[-1]
done = False
steps = 0
while steps < 200:
    action, prob, val, h_next = choose_action_value(
        obs, h_prev, actor, critic)
    next_observation, reward, done, info = env.step([action])
    observation_list.append(observation)
    observation = next_observation
    obs = observation[-1]
    steps += 1
    h_prev = h_next
    if done:
        break

return total_loss_list, observation_list, loss_list, reward_list

```

Clipped and GAE

```

In [10]: N = 100 # number of times to collect new trajectories and update actor
M = 1 # num of trajectories
max_trajectory_len = 200 # trajectory length
batch_size = 25 # size for minibatch
n_epochs = 30 # number of epochs to optimize loss
alpha = 0.0001
hidden_layers = 2
hidden_dims = 64
gamma = 0.7
lmbda = 0.9
clip_value = 0.1
use_gae = True
use_clip = True

total_losses, observations, loss_list, reward_list = ppo(
    N, M, max_trajectory_len, batch_size, alpha, hidden_layers, hidden_dims,
    gamma, lmbda, clip_value, use_gae, use_clip)

```

D:\miniconda3\envs\cmpe260\lib\site-packages\gym\utils\passive_env_checker.py:195: UserWarning: WARN: The result returned by `env.reset()` was not a tuple of the form `(obs, info)`, where `obs` is a observation and `info` is a dictionary containing additional information. Actual type: `
 logger.warn(
D:\miniconda3\envs\cmpe260\lib\site-packages\gym\utils\passive_env_checker.py:219: DeprecationWarning: WARN: Core environment is written in old step API which returns one bool instead of two. It is recommended to rewrite the environment with new step API.
 logger.deprecation(

```

In [11]: tmp = [x for x in range(100)]


```

```

In [12]: plt.plot(tmp, np.array([l.cpu().detach() for l in loss_list]))


```

Out[12]: [



```
In [13]: plt.plot(tmp, np.array([l.cpu().detach() for l in reward_list]))
```

Out[13]: [



```
In [14]: o = []
for i in range(N):
    o.append(observations[200*(i+1)-200:200*(i+1)])
```

```
In [15]: o_old = o.copy()
o = [j for i, j in enumerate(o) if i%10==0]
```

```
In [16]: o = o_old[-10:]
```

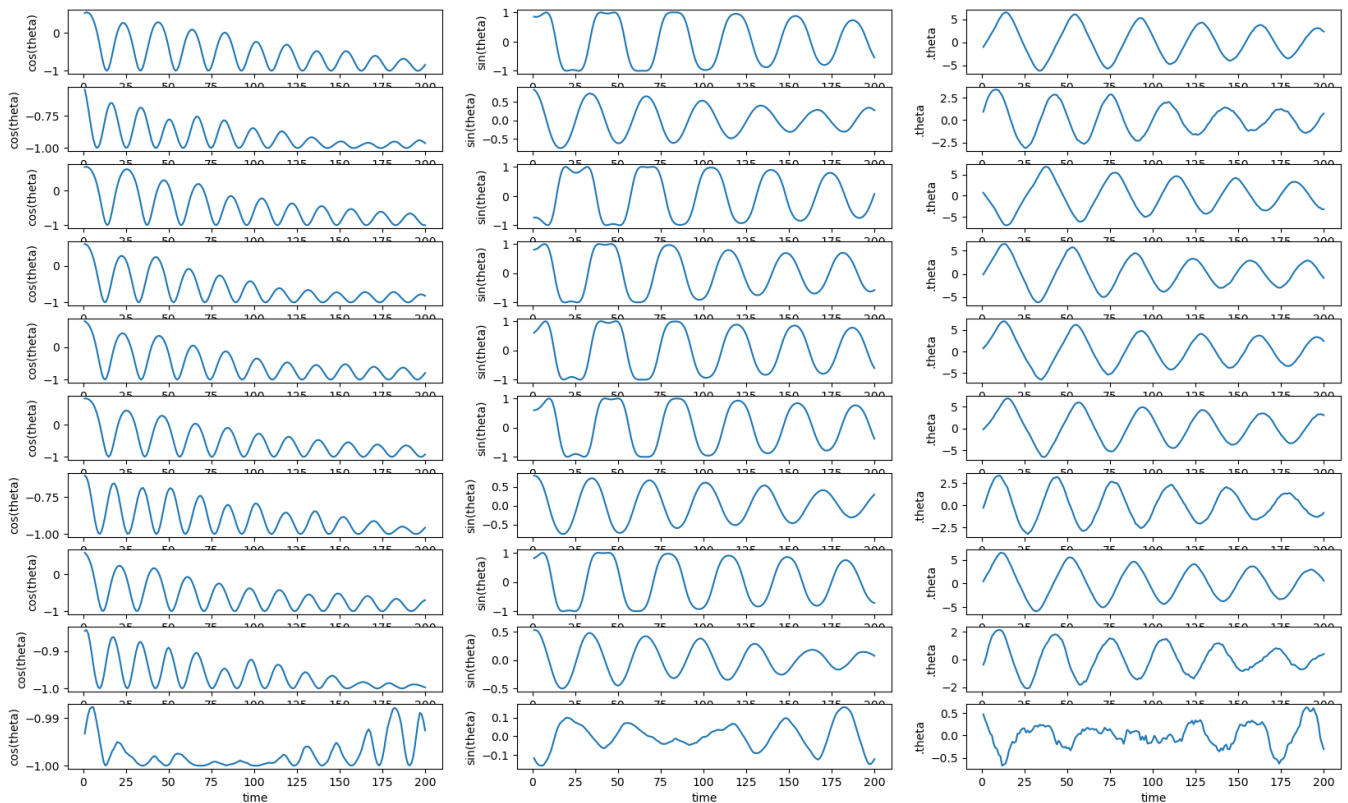


```
In [17]: x_ = [i+1 for i in range(200)]
```

```
In [18]: np.array(o).shape
```

```
Out[18]: (10, 200, 3)
```

```
In [19]: N = 10
plt.figure(1, figsize=(20, 12))
for i in range(N):
    obs = o[i]
    # print(obs)
    theta_ = []
    x_list = []
    y_list = []
    for x,y,v in obs:
        # print(x, y)
        theta_.append(v)
        x_list.append(x)
        y_list.append(y)
    plt.subplot(N, 3, 3*i+1)
    plt.plot(x_, x_list)
    plt.xlabel('time')
    plt.ylabel('cos(theta)')
    plt.subplot(N, 3, 3*i+2)
    plt.plot(x_, y_list)
    plt.xlabel('time')
    plt.ylabel('sin(theta)')
    plt.subplot(N, 3, 3*i+3)
    plt.plot(x_, theta_)
    plt.xlabel('time')
    plt.ylabel('theta')
```




```
In [20]: len(total_losses)
```

```
Out[20]: 24000
```

```
In [21]: x = np.arange(len(total_losses))
y = [float(x) for x in total_losses]
plt.figure(figsize=(8, 6))
plt.plot(x, y)
```

Out[21]: [



Clipped and No GAE


```
In [22]: N = 100 # number of times to collect new trajectories and update actor
M = 1 # num of trajectories
max_trajectory_len = 200 # trajectory length
batch_size = 25 # size for minibatch
n_epochs = 30 # number of epochs to optimize loss
alpha = 0.0001
hidden_layers = 2
hidden_dims = 64
gamma = 0.7
lmbda = 0.9
clip_value = 0.1
use_gae = False
use_clip = True

total_losses, observations, loss_list, reward_list = ppo(
    N, M, max_trajectory_len, batch_size, alpha, hidden_layers, hidden_dims,
    gamma, lmbda, clip_value, use_gae, use_clip)
```

```
In [23]: tmp = [x for x in range(100)]
```


```
In [24]: plt.plot(tmp, np.array([l.cpu().detach() for l in loss_list]))
```

Out[24]: [



```
In [25]: plt.plot(tmp, np.array([l.cpu().detach() for l in reward_list]))
```

```
Out[25]: [matplotlib.lines.Line2D at 0x251399c6050]
```



```
In [26]: o = []
for i in range(N):
    o.append(observations[200*(i+1)-200:200*(i+1)])
```

```
In [27]: o_old = o.copy()
o = [j for i, j in enumerate(o) if i%10==0]
```

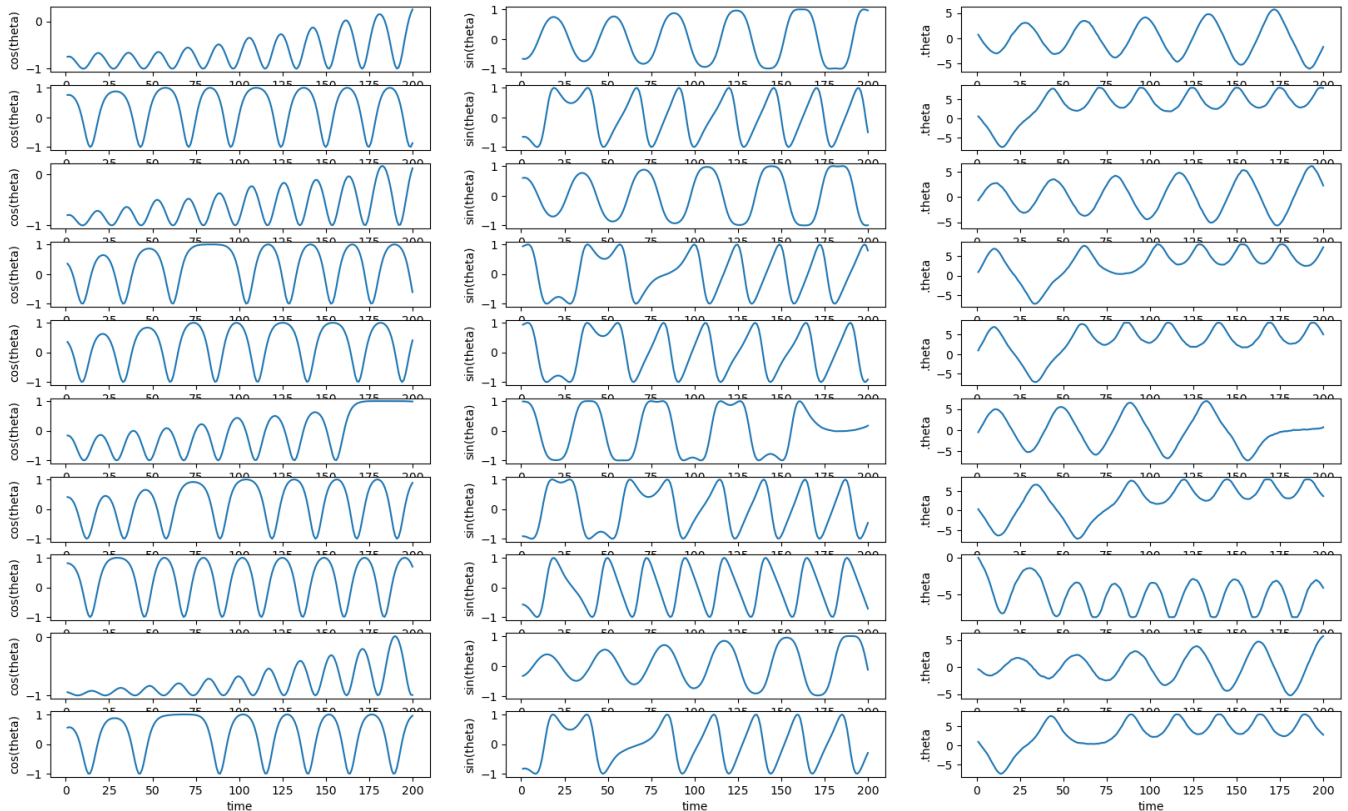
```
In [28]: o = o_old[-10:]
```

```
In [29]: x_ = [i+1 for i in range(200)]
```

```
In [30]: np.array(o).shape
```

```
Out[30]: (10, 200, 3)
```

```
In [31]: N = 10
plt.figure(1, figsize=(20, 12))
for i in range(N):
    obs = o[i]
    # print(obs)
    theta_ = []
    x_list = []
    y_list = []
    for x,y,v in obs:
        # print(x, y)
        theta_.append(v)
        x_list.append(x)
        y_list.append(y)
    plt.subplot(N, 3, 3*i+1)
    plt.plot(x_, x_list)
    plt.xlabel('time')
    plt.ylabel('cos(theta)')
    plt.subplot(N, 3, 3*i+2)
    plt.plot(x_, y_list)
    plt.xlabel('time')
    plt.ylabel('sin(theta)')
    plt.subplot(N, 3, 3*i+3)
    plt.plot(x_, theta_)
    plt.xlabel('time')
    plt.ylabel('theta')
```




```
In [32]: len(total_losses)
```

```
Out[32]: 24000
```

```
In [33]: x = np.arange(len(total_losses))
y = [float(x) for x in total_losses]
plt.figure(figsize=(8, 6))
plt.plot(x, y)
```

Out[33]: [



No Clipping and GAE


```
In [34]: N = 100 # number of times to collect new trajectories and update actor
M = 1 # num of trajectories
max_trajectory_len = 200 # trajectory length
batch_size = 25 # size for minibatch
n_epochs = 30 # number of epochs to optimize loss
alpha = 0.0001
hidden_layers = 2
hidden_dims = 64
gamma = 0.7
lmbda = 0.9
clip_value = 0.1
use_gae = True
use_clip = False

total_losses, observations, loss_list, reward_list = ppo(
    N, M, max_trajectory_len, batch_size, alpha, hidden_layers, hidden_dims,
    gamma, lmbda, clip_value, use_gae, use_clip)
```

```
In [35]: tmp = [x for x in range(100)]
```


```
In [36]: plt.plot(tmp, np.array([l.cpu().detach() for l in loss_list]))
```

Out[36]: [



```
In [37]: plt.plot(tmp, np.array([l.cpu().detach() for l in reward_list]))
```

```
Out[37]: [matplotlib.lines.Line2D at 0x251107b5210]
```



```
In [38]: o = []
for i in range(N):
    o.append(observations[200*(i+1)-200:200*(i+1)])
```

```
In [39]: o_old = o.copy()
o = [j for i, j in enumerate(o) if i%10==0]
```

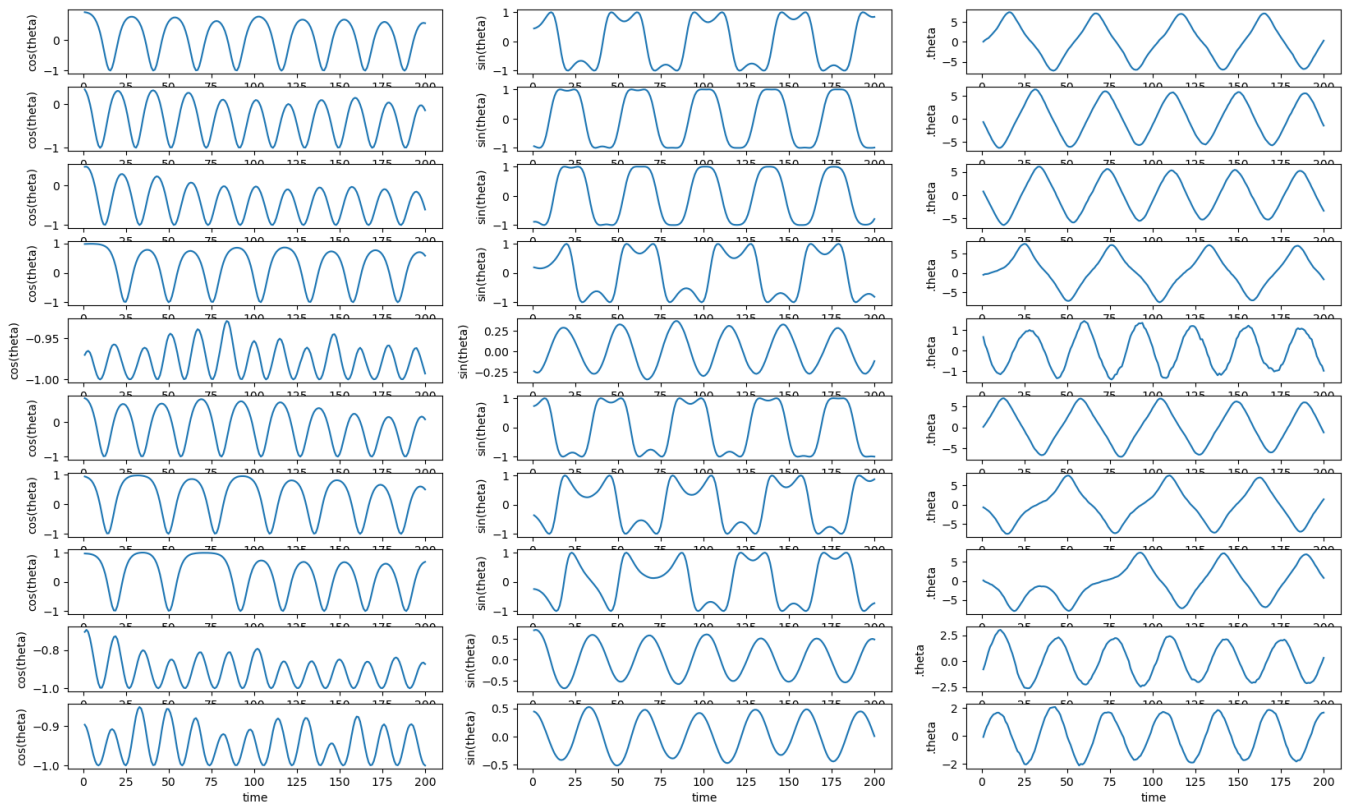
```
In [40]: o = o_old[-10:]
```

```
In [41]: x_ = [i+1 for i in range(200)]
```

```
In [42]: np.array(o).shape
```

```
Out[42]: (10, 200, 3)
```

```
In [43]: N = 10
plt.figure(1, figsize=(20, 12))
for i in range(N):
    obs = o[i]
    # print(obs)
    theta_ = []
    x_list = []
    y_list = []
    for x,y,v in obs:
        # print(x, y)
        theta_.append(v)
        x_list.append(x)
        y_list.append(y)
    plt.subplot(N, 3, 3*i+1)
    plt.plot(x_, x_list)
    plt.xlabel('time')
    plt.ylabel('cos(theta)')
    plt.subplot(N, 3, 3*i+2)
    plt.plot(x_, y_list)
    plt.xlabel('time')
    plt.ylabel('sin(theta)')
    plt.subplot(N, 3, 3*i+3)
    plt.plot(x_, theta_)
    plt.xlabel('time')
    plt.ylabel('theta')
```




```
In [44]: len(total_losses)
```

```
Out[44]: 24000
```

```
In [45]: x = np.arange(len(total_losses))
y = [float(x) for x in total_losses]
plt.figure(figsize=(8, 6))
plt.plot(x, y)
```

```
Out[45]: [<matplotlib.lines.Line2D at 0x25136191120>]
```



No Clipping and No GAE


```
In [46]: N = 100 # number of times to collect new trajectories and update actor
M = 1 # num of trajectories
max_trajectory_len = 200 # trajectory length
batch_size = 25 # size for minibatch
n_epochs = 30 # number of epochs to optimize loss
alpha = 0.0001
hidden_layers = 2
hidden_dims = 64
gamma = 0.7
lmbda = 0.9
clip_value = 0.1
use_gae = False
use_clip = False

total_losses, observations, loss_list, reward_list = ppo(
    N, M, max_trajectory_len, batch_size, alpha, hidden_layers, hidden_dims,
    gamma, lmbda, clip_value, use_gae, use_clip)
```

```
In [47]: tmp = [x for x in range(100)]
```


```
In [48]: plt.plot(tmp, np.array([l.cpu().detach() for l in loss_list]))
```

```
Out[48]: [ <matplotlib.lines.Line2D at 0x250d40ce7a0> ]
```

```
In [49]: plt.plot(tmp, np.array([l.cpu().detach() for l in reward_list]))
```

```
Out[49]: [matplotlib.lines.Line2D at 0x250d4129780]
```



```
In [50]: o = []
for i in range(N):
    o.append(observations[200*(i+1)-200:200*(i+1)])
```

```
In [51]: o_old = o.copy()
o = [j for i, j in enumerate(o) if i%10==0]
```

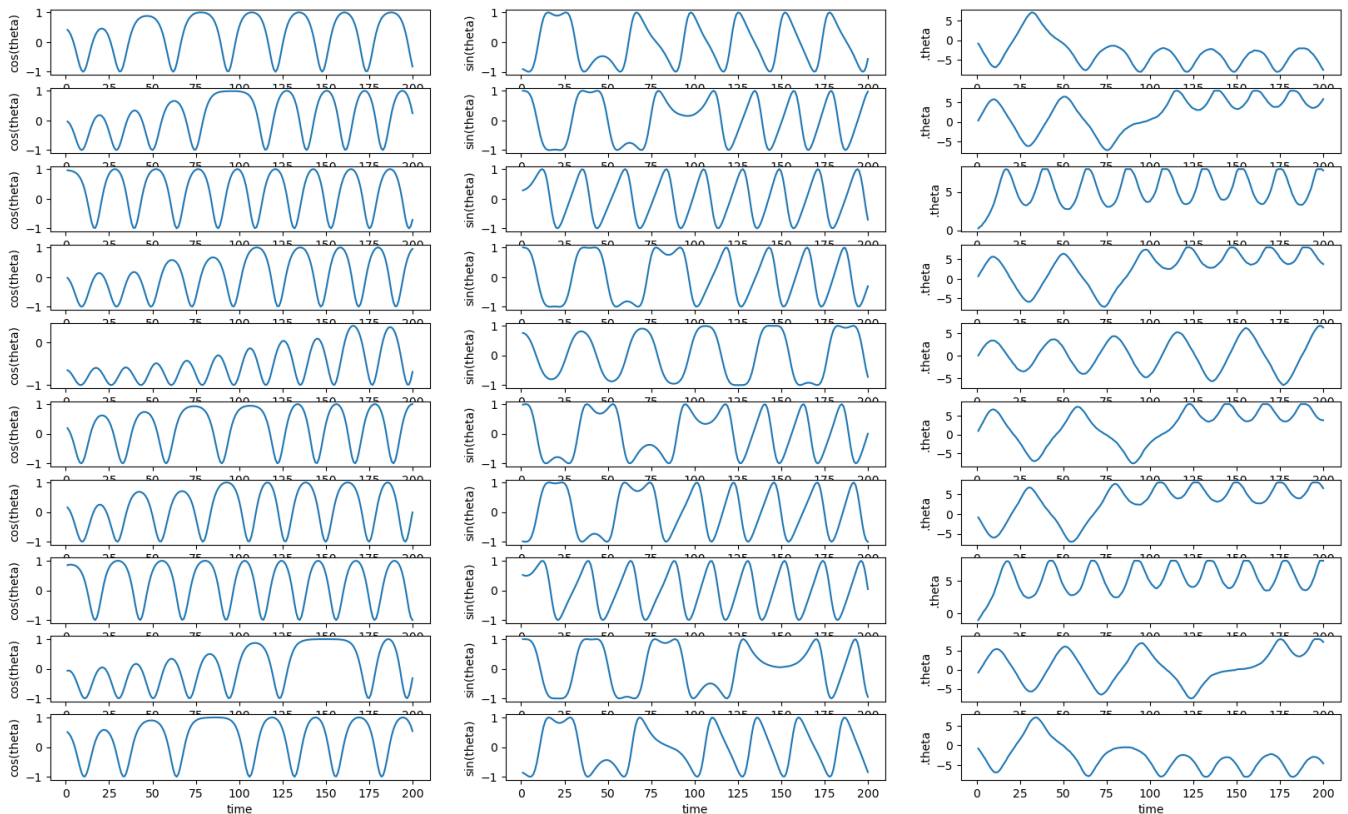
```
In [52]: o = o_old[-10:]
```

```
In [53]: x_ = [i+1 for i in range(200)]
```

```
In [54]: np.array(o).shape
```

```
Out[54]: (10, 200, 3)
```

```
In [55]: N = 10
plt.figure(1, figsize=(20, 12))
for i in range(N):
    obs = o[i]
    # print(obs)
    theta_ = []
    x_list = []
    y_list = []
    for x,y,v in obs:
        # print(x, y)
        theta_.append(v)
        x_list.append(x)
        y_list.append(y)
    plt.subplot(N, 3, 3*i+1)
    plt.plot(x_, x_list)
    plt.xlabel('time')
    plt.ylabel('cos(theta)')
    plt.subplot(N, 3, 3*i+2)
    plt.plot(x_, y_list)
    plt.xlabel('time')
    plt.ylabel('sin(theta)')
    plt.subplot(N, 3, 3*i+3)
    plt.plot(x_, theta_)
    plt.xlabel('time')
    plt.ylabel('theta')
```



```
In [56]: len(total_losses)
```

```
Out[56]: 24000
```

```
In [57]: x = np.arange(len(total_losses))
y = [float(x) for x in total_losses]
plt.figure(figsize=(8, 6))
plt.plot(x, y)
```

Out[57]: [

